

Arttu Kuitunen
Student number 1500155
February, 2, 2025

Exercise 1 | TKO_7092 Evaluation of Machine Learning Methods 2025

Prediction of the metal ion content from multi-parameter data

Use K-Nearest Neighbor Regression with euclidean distance to predict total metal concentration (c_total), concentration of Cadmium (Cd) and concentration of Lead (Pb), using number of neighbors k = 1, 3, 5, 7.

- You may use Nearest Neighbor Regression from <https://scikit-learn.org/stable/modules/neighbors.html>
- Implement Leave-One-Out cross-validation and calculate the C-index for each output (c_total, Cd, Pb).
- Implement Leave-Replicas-Out cross-validation and calculate the C-index for each output (c_total, Cd, Pb).
- Return your solution as a Jupyter Notebook .ipynb notebook and as a PDF-file made from it. Please, add your full name to the file name.
- The exercise will be graded by a maximum of 2 points.
- Submit to moodle your solution on **** Wednesday 5 of February **** at the latest.

Please be aware that you are required to submit your individual solution.
Submissions with identical or similar code will result in a failure for the exercise.

Import libraries

```
In [199... # In this cell import all libraries you need. For example:
import numpy as np
import pandas as pd
from scipy.stats import somersd
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import LeaveOneOut
import matplotlib.pyplot as plt
```

Read and visualize the dataset

Note: This dataset differs slightly from the one used in the video lectures.

In this dataset, some mixtures have 3 replicas, while others have 4 replicas.

In the following cell:

- Read the file water_data.csv
- Print the dimensions of the dataset (i.e. number of rows and columns) and display the first 5 rows.
- Identify the inputs and the outputs columns.
- Provide the number of mixtures with 3 replicas and 4 replicas, respectively.

```
In [200... # Maybe I want average of the replicas, since some have 3 and some 4 replicas
data = pd.read_csv('water_data.csv')
input_columns = ['Mod1', 'Mod2', 'Mod3']
output_columns = ['c_total', 'Cd', 'Pb']

outputs = data[input_columns]
inputs = data[output_columns]

print("Shape (rows, columns):\n", data.shape),
print("\nFirst 5 rows:\n", data.head()),
print("Inputs:", inputs.columns),
print("Outputs:", outputs.columns)
```

Shape (rows, columns):
(225, 6)

First 5 rows:

	c_total	Cd	Pb	Mod1	Mod2	Mod3
0	0.0	0.0	0.0	-0.999216	-0.714208	-0.414911
1	0.0	0.0	0.0	-0.990800	-0.714373	-0.238335
2	0.0	0.0	0.0	-0.990539	-0.714125	0.020788
3	14.0	0.0	14.0	-1.001247	-0.713546	0.945465
4	14.0	0.0	14.0	-1.013727	-0.714125	0.569631

Inputs: Index(['c_total', 'Cd', 'Pb'], dtype='object')

Outputs: Index(['Mod1', 'Mod2', 'Mod3'], dtype='object')

```
In [201... # The amount of occurrences in dataset
four = (data.groupby(output_columns).size() == 4).sum()
three = (data.groupby(output_columns).size() == 3).sum()

if four * 4 + three * 3 == len(data):
    print("Matching amount of replicas to data")
else:
    print("Not matching amount of replicas to data")
```

Matching amount of replicas to data

C-index code

```
In [202... # In this cell is the fuction that computes the c-index value based on Somers'D sta
# Use this fuction as the evaluation metric in the Leave-One-Out (LOOCV) and Leave-

def cindex(true, pred):
    s_d = somersd(true, y=pred, alternative='two-sided')
```

```
c_index = (s_d.statistic + 1.0)/2.0
return c_index
```

Functions

```
In [203... # In this cell add the functions that you need for the data analysis part.

# Leave-one-out cross-validation
def leave_one_out_cv(df, k_values):
    loo = LeaveOneOut()
    results = {k: {col: [] for col in output_columns} for k in k_values}

    # Iterating every train-test split possible
    for train_idx, test_idx in loo.split(df):
        train_data, test_data = df.iloc[train_idx], df.iloc[test_idx]
        X_train, y_train = train_data[input_columns], train_data[output_columns]
        X_test, y_test = test_data[input_columns], test_data[output_columns]

        # Fit models for every k neighbors
        for k in k_values:
            model = KNeighborsRegressor(n_neighbors=k, metric='euclidean')
            model.fit(X_train, y_train)
            predictions = model.predict(X_test)

            # set results for outputs
            for i, col in enumerate(output_columns):
                results[k][col].append((y_test.iloc[0][col], predictions[0][i]))

    c_index_scores = {k: {col: cindex(*zip(*results[k][col])) for col in output_col}
                      for k in k_values}
    return c_index_scores

# Leave-replicas-out cross-validation
def leave_replicas_out_cv(df, k_values):
    results = {k: {col: [] for col in output_columns} for k in k_values}

    # Get all unique groups
    groups = list(df.groupby(output_columns).groups.values())

    # For each sample in each group
    for test_indices in zip(*groups): # Combines one replica from each group

        # Create combined train/test split
        test_data = df.loc[list(test_indices)]
        train_data = df.drop(list(test_indices), axis=0)

        X_train, y_train = train_data[input_columns], train_data[output_columns]
        X_test, y_test = test_data[input_columns], test_data[output_columns]

        # Fit models for every k neighbors
        for k in k_values:
            model = KNeighborsRegressor(n_neighbors=k, metric='euclidean')
            model.fit(X_train, y_train)
            predictions = model.predict(X_test)
```

```

        for i, col in enumerate(output_columns):
            for j in range(len(test_indices)):
                results[k][col].append((y_test.iloc[j][col], predictions[j][i]))

    c_index_scores = {k: {col: cindex(*zip(*results[k][col])) for col in output_col
    return c_index_scores

```

In [204... *# Some list comprehension help I found.*
<https://stackoverflow.com/questions/40646458/list-comprehension-in-pandas>

Cross validation practical python example
<https://www.youtube.com/watch?v=-8s9KuNo5SA>

Results for Leave-One-Out cross-validation

In [205... *# Here run your script for Leave-One-Out cross-validation and print the correspondi*

```

k_values = [1, 3, 5, 7]
loocv_results = leave_one_out_cv(data, k_values)
print("\nLeave-one-out Cross-validation results:")
for k, scores in loocv_results.items():
    print(f"k={k}: {scores}")

```

Leave-one-out Cross-validation results:

```

k=1: {'c_total': 0.9082833811137173, 'Cd': 0.921869127656909, 'Pb': 0.88054871173842
23}
k=3: {'c_total': 0.9141907740422205, 'Cd': 0.8995907629348143, 'Pb': 0.8744519146448
406}
k=5: {'c_total': 0.8941012944140387, 'Cd': 0.8619660082682591, 'Pb': 0.8542614941328
768}
k=7: {'c_total': 0.8737294761532447, 'Cd': 0.8141520858562659, 'Pb': 0.8355326345680
043}

```

Results for Leave-Replicas-Out cross-validation

In [206... *# Here run your script for Leave-Replicas-Out cross-validation and print the corres*

```

lrocvc_results = leave_replicas_out_cv(data, k_values)
print("\nLeave-replicas-out cross-validation results:")
for k, scores in lrocvc_results.items():
    print(f"k={k}: {scores}")

```

Leave-replicas-out cross-validation results:

```
k=1: {'c_total': 0.9147659389594873, 'Cd': 0.9208777987026575, 'Pb': 0.888653483992467}
k=3: {'c_total': 0.9218529379819702, 'Cd': 0.8964741577735929, 'Pb': 0.8835268884703913}
k=5: {'c_total': 0.8951069838166612, 'Cd': 0.8438480853735091, 'Pb': 0.8612680477087257}
k=7: {'c_total': 0.8794395568589117, 'Cd': 0.7924513496547395, 'Pb': 0.8411801632140615}
```

```
k=1: {'c_total': 0.9147659389594873, 'Cd': 0.9208777987026575, 'Pb': 0.888653483992467}
k=3: {'c_total': 0.9218529379819702, 'Cd': 0.8964741577735929, 'Pb': 0.8835268884703913}
k=5: {'c_total': 0.8951069838166612, 'Cd': 0.8438480853735091, 'Pb': 0.8612680477087257}
k=7: {'c_total': 0.8794395568589117, 'Cd': 0.7924513496547395, 'Pb': 0.8411801632140615}
```

Plot Leave-One-Out and Leave-Replicas-Out Results

Note: You may plot the results as they were presented in the video lecture (refer to MOOC2-Module 2 .pptx slides).

In [207...

```
# Create figure with 3 subplots

# Axes for each output
fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(15, 5))
k_values = [1, 3, 5, 7]

loocv_plot = {col: [loocv_results[k][col] for k in k_values] for col in output_columns}
lrocvcv_plot = {col: [lrocvcv_results[k][col] for k in k_values] for col in output_columns}

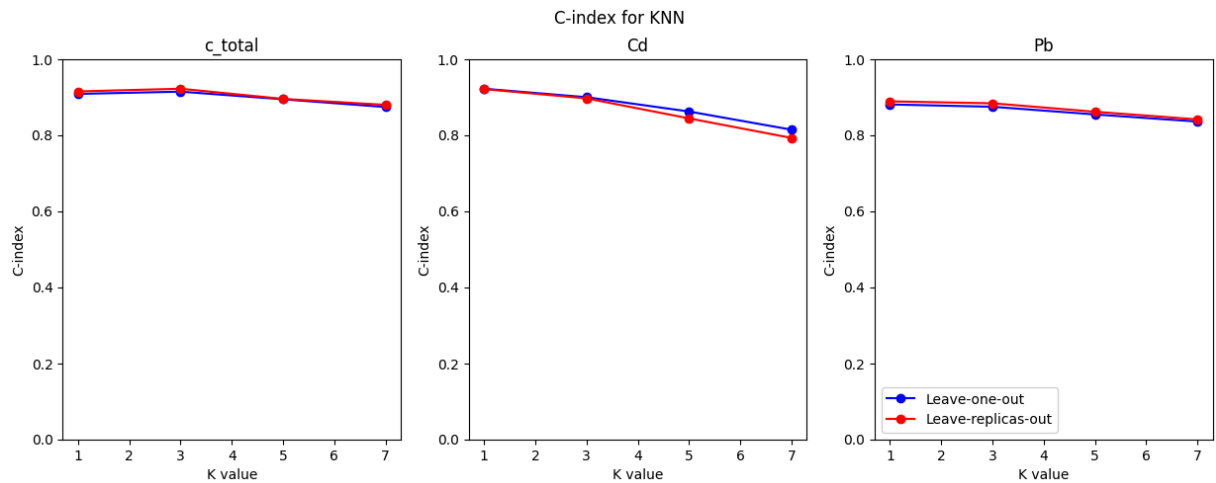
# Plot with lines connecting points
ax1.plot(k_values, loocv_plot['c_total'], 'bo-')
ax1.plot(k_values, lrocvcv_plot['c_total'], 'ro-')

ax2.plot(k_values, loocv_plot['Cd'], 'bo-')
ax2.plot(k_values, lrocvcv_plot['Cd'], 'ro-')

ax3.plot(k_values, loocv_plot['Pb'], 'bo-')
ax3.plot(k_values, lrocvcv_plot['Pb'], 'ro-')

# Set titles and labels
titles = ['c_total', 'Cd', 'Pb']
fig.suptitle('C-index for KNN')
for ax, title in zip([ax1, ax2, ax3], titles):
    ax.set_title(title)
    ax.set_xlabel('K value')
    ax.set_ylabel('C-index')
    ax.set_ylim(0, 1)

plt.legend(['Leave-one-out', 'Leave-replicas-out'])
plt.show()
```



Interpretation of results

Answer the following questions based on the results obtained

1. Which cross-validation approach produced more optimistic results, and why?
2. Which cross-validation method provides a better estimate of the model's performance on unseen mixtures? Explain your answer.

Answers:

1. At first, leave-one-out was highly more optimal, showing similiar difference than in the lecture slideshow. At first I performed leave-replicas-out in more of leave-group-out fashion, where every group was trained individually, which in my opinion was not the goal here. Leave-one-out has more training data, therefore it should perform better. It takes more computing.
2. I guess leave-replicas-out, since it is less likely to over fitting.